

— Linear Regression

Kinds of ML

- **Supervised learning** - Given X , can we predict Y ?
- **Unsupervised learning** - What patterns can we find in X ?
There is no Y .
- **Reinforcement learning** - How can a virtual agent maximize a reward.



Supervised learning

- **Classification**
- **Regression**

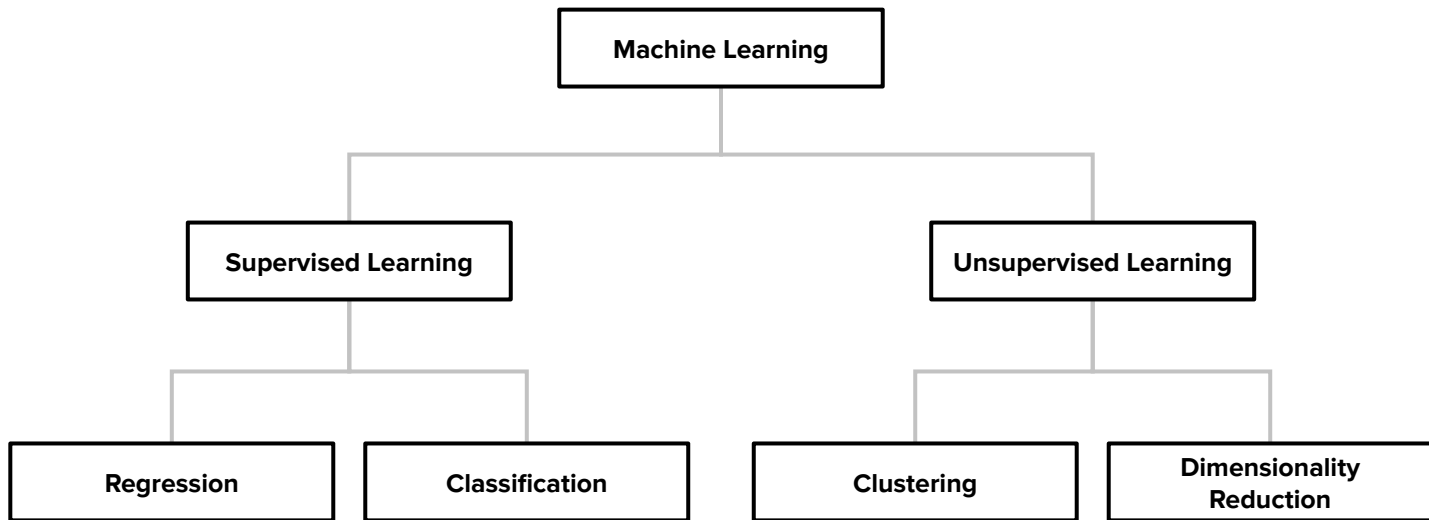


Two Kinds of Supervised Learning


Regression: y -variable is a continuous number

- *Crop yields*
- *House sale price*
- *Other examples?*


Roadmap of ML




Roadmap of ML



Roadmap of ML



Roadmap of ML



Supervised Learning Transparency

Data



[img source](#)

???

Profitable Results!




Supervised Learning Transparency

Data



[img source](#)



Profitable Results!



Let's do it!



— Linear Regression

Supervised, white-box, regression



Ordinary least squares linear regression (**OLS**)

Predict response variable (y) from at least one independent variable (x).


$$y = \beta_0 + \beta_1 x + \varepsilon$$

Ordinary least squares linear regression (**OLS**)


$$y = \beta_0 + \beta_1 x + \varepsilon$$

The game is “find the best betas.”


Graphically, this is:



Graphically, this is:




Graphically, this is:



Graphically, this is:

The difference between the actual and the predicted value is called a **residual**.


The line of “best fit” **minimizes the sum of the squared residuals**.



Graphically, this is:

Minimize the **sum of the squared residuals**

OLS: “ordinary least squares”



The Residual

A **residual** is the same as an error

$$e_i = y_i - \hat{y}_i$$


This measures how “off” a prediction was.

Sum of Squared Errors (SSE)

$$SSE = \sum (y_i - \hat{y}_i)^2 = \sum e_i^2$$

Fitting OLS models

This is the quantity we seek to **minimize** to find the best values for our betas!



Multiple Linear Regression

Supervised, white-box, regression



Multiple Linear Regression

Generally, you'll have more than one predictor variable.

Formula for a predicted y value:

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 \mathbf{x}_1 + \hat{\beta}_2 \mathbf{x}_2 + \dots \hat{\beta}_p \mathbf{x}_p + \varepsilon$$

Notation: Design Matrix

Mathematically speaking, every time we fit a model, we need a data matrix, sometimes called a **design matrix**:

$$X = \begin{bmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1p} \\ 1 & x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_{n1} & x_{n2} & \cdots & x_{np} \end{bmatrix}$$

Notation: Design Matrix

$$X = \begin{bmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1p} \\ 1 & x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_{n1} & x_{n2} & \cdots & x_{np} \end{bmatrix}$$

Each **row** is an **observation**.

Notation: Design Matrix

$$X = \begin{bmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1p} \\ 1 & x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_{n1} & x_{n2} & \cdots & x_{np} \end{bmatrix}$$

Each **column**
is a variable.

Notation: Design Matrix

$$X = \begin{bmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1p} \\ 1 & x_{21} & x_{22} & \cdots & x_{2p} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_{n1} & x_{n2} & \cdots & x_{np} \end{bmatrix}$$

The first column is all 1s and corresponds to the **intercept**. (scikit-learn handles this automatically)

Multiple linear regression

Interpreting results

Now when you interpret a beta weight, you must use the caveat: “holding everything else constant”.

Multiple linear regression

Interpreting results example:

If β_1 is 5.0 then a 1 degree celsius increase in temperature is associated with a 5MW increase in energy usage, **holding everything else constant.**

Assumption

For inference to make sense:

There is no **M**ulticollinearity in the predictor variables.

Check for highly correlated variables using `df.corr` & `seaborn`. (Or better, `VIF` in `statsmodels`. Picks up patterns of correlation.).

Let's do it!



The Residual

A **residual** is:

$$e_i = y_i - \hat{y}_i$$

How far “off” a prediction was.


The error.

Sum of Squared Errors (SSE)

$$SSE = \sum (y_i - \hat{y}_i)^2 = \sum e_i^2$$

Fitting OLS models

This is the quantity we seek to **minimize** to find the best values for our betas!



— LINE Assumptions

OLS Assumptions


OLS comes with some assumptions. Some matter for performance and some matter for inference.

OLS Regression Assumptions

- **L** - Linearity. Relationship between x and y should be approximately linear.
- **I** - Independence of errors. (Time series issue)
- **N** - Normality. Residuals should be approximately normally distributed.
- **E** - Equal variances, aka “**homoscedasticity**”. Residuals should have approximately equal variances.



L is for Linearity



L is for Linearity

We will learn how to transform X or y to try to make the relationship more linear.

Or maybe another algorithm would be better.




I is for Independence

Are our samples independent from one another?


With time series, they usually are not independent. The previous row tells us something about the next row.

N is for Normality of errors



E is for Equal Variances


We can see a parabolic pattern in our residual plot.



E is for Equal Variance

We want to see **absolute randomness** in our residual plots. That is, no pattern whatsoever.

Here's an example of an ideal residual plot.



What to do if our **LINE** assumptions are violated?

A common scenario:

- A slightly **curvilinear** relationship between x and y
- Very right-skew residuals
- Residuals that tend to spread out from right to left (a “fan shape”)

A log transformation often helps fix these issues

LINEM Assumptions

- **L** - Linear relationship
- **I** - Independent errors (time series issue, often)
- **N** - Normally distributed errors
- **E** - Equal variance of errors (homoscedasticity)

Add an M for multiple regression (more than 1 X)

- **M** - No multicollinearity (highly correlated predictors)



Let's do it!

