

— Recurrent Neural Networks (RNNs)

Agenda



Introduction to RNNs



Case Study: Using RNNs to predict stock prices



RNN in Python

Learning Objectives

- Preprocess sequence data for RNN modeling
- Design, train and evaluate an RNN model
- Create a model that predicts the price of Apple's (AAPL) stock



RNNs



Introduction to RNNs



Why do we need RNNs?

Let's start with an example. Suppose I asked you to tell me which direction this ball would move next?


Why do we need RNNs?

Let's start with an example. Suppose I asked you to tell me which direction this ball would move next?




Why do we need RNNs?

Let's start with an example. Suppose I asked you to tell me which direction this ball would move next?




Why do we need RNNs?

Let's start with an example. Suppose I asked you to tell me which direction this ball would move next?



Why do we need RNNs?

Let's start with an example. Suppose I asked you to tell me which direction this ball would move next?



Sequence Data

RNNs are used on sequential data.

What are some examples of sequential data?

Sequence Data


RNNs are used on sequential data.

What are some examples of sequential data?

- Language/Text
- Timeseries
- Sound




Reminder: Standard (Feed Forward) Neural Networks



$$\mathbf{z} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

$$\mathbf{y} = g(\mathbf{z})$$

RNNs



$$z = \mathbf{W}_H \mathbf{h}_{t-1} + \mathbf{W}_x \mathbf{x} + \mathbf{b}$$
$$\mathbf{y} = g(\mathbf{z})$$

RNNs



Case Study: using RNNs to predict stock prices



Keras TimeseriesGenerator

	Volume	Open	Close	Target
2019-01-01	1,000,000	\$150.00	\$155.00	NaN
2019-01-02	2,000,000	\$160.00	\$165.00	1
2019-01-03	3,000,000	\$155.00	\$160.00	0
2019-01-04	4,000,000	\$180.00	\$185.00	1
2019-01-05	5,000,000	\$175.00	\$180.00	0
2019-01-06	6,000,000	\$170.00	\$175.00	0

Keras TimeseriesGenerator: Length = 3

	Volume	Open	Close	Target
2019-01-01	1,000,000	\$150.00	\$155.00	NaN
2019-01-02	2,000,000	\$160.00	\$165.00	1
2019-01-03	3,000,000	\$155.00	\$160.00	0
2019-01-04	4,000,000	\$180.00	\$185.00	1
2019-01-05	5,000,000	\$175.00	\$180.00	0
2019-01-06	6,000,000	\$170.00	\$175.00	0

Keras TimeseriesGenerator: Length = 3

	Volume	Open	Close	Target
2019-01-01	1,000,000	\$150.00	\$155.00	NaN
2019-01-02	2,000,000	\$160.00	\$165.00	1
2019-01-03	3,000,000	\$155.00	\$160.00	0
2019-01-04	4,000,000	\$180.00	\$185.00	1
2019-01-05	5,000,000	\$175.00	\$180.00	0
2019-01-06	6,000,000	\$170.00	\$175.00	0

Keras TimeseriesGenerator: Length = 3


	Volume	Open	Close	Target
2019-01-01	1,000,000	\$150.00	\$155.00	NaN
2019-01-02	2,000,000	\$160.00	\$165.00	1
2019-01-03	3,000,000	\$155.00	\$160.00	0
2019-01-04	4,000,000	\$180.00	\$185.00	1
2019-01-05	5,000,000	\$175.00	\$180.00	0
2019-01-06	6,000,000	\$170.00	\$175.00	0

Keras TimeseriesGenerator: Length = 4

	Volume	Open	Close	Target
2019-01-01	1,000,000	\$150.00	\$155.00	NaN
2019-01-02	2,000,000	\$160.00	\$165.00	1
2019-01-03	3,000,000	\$155.00	\$160.00	0
2019-01-04	4,000,000	\$180.00	\$185.00	1
2019-01-05	5,000,000	\$175.00	\$180.00	0
2019-01-06	6,000,000	\$170.00	\$175.00	0


RNN Architecture

	Volume	Open	Close
2019-01-01	1,000,000	\$150.00	\$155.00
2019-01-02	2,000,000	\$160.00	\$165.00
2019-01-03	3,000,000	\$155.00	\$160.00




RNN Architecture

	Volume	Open	Close
2019-01-01	1,000,000	\$150.00	\$155.00
2019-01-02	2,000,000	\$160.00	\$165.00
2019-01-03	3,000,000	\$155.00	\$160.00




RNN Architecture

	Volume	Open	Close
2019-01-01	1,000,000	\$150.00	\$155.00
2019-01-02	2,000,000	\$160.00	\$165.00
2019-01-03	3,000,000	\$155.00	\$160.00




RNN Architecture

	Volume	Open	Close
2019-01-01	1,000,000	\$150.00	\$155.00
2019-01-02	2,000,000	\$160.00	\$165.00
2019-01-03	3,000,000	\$155.00	\$160.00




RNN Architecture

	Volume	Open	Close
2019-01-01	1,000,000	\$150.00	\$155.00
2019-01-02	2,000,000	\$160.00	\$165.00
2019-01-03	3,000,000	\$155.00	\$160.00




RNN Architecture

	Volume	Open	Close
2019-01-01	1,000,000	\$150.00	\$155.00
2019-01-02	2,000,000	\$160.00	\$165.00
2019-01-03	3,000,000	\$155.00	\$160.00




RNN Architecture: Unfolded

	Volume	Open	Close
2019-01-01	1,000,000	\$150.00	\$155.00
2019-01-02	2,000,000	\$160.00	\$165.00
2019-01-03	3,000,000	\$155.00	\$160.00




RNN Architecture: Unfolded

	Volume	Open	Close
2019-01-01	1,000,000	\$150.00	\$155.00
2019-01-02	2,000,000	\$160.00	\$165.00
2019-01-03	3,000,000	\$155.00	\$160.00




RNN Architecture: Unfolded

	Volume	Open	Close
2019-01-01	1,000,000	\$150.00	\$155.00
2019-01-02	2,000,000	\$160.00	\$165.00
2019-01-03	3,000,000	\$155.00	\$160.00




RNN Architecture: Unfolded

	Volume	Open	Close
2019-01-01	1,000,000	\$150.00	\$155.00
2019-01-02	2,000,000	\$160.00	\$165.00
2019-01-03	3,000,000	\$155.00	\$160.00




RNN Architecture: Unfolded

	Volume	Open	Close
2019-01-01	1,000,000	\$150.00	\$155.00
2019-01-02	2,000,000	\$160.00	\$165.00
2019-01-03	3,000,000	\$155.00	\$160.00



RNN Architecture: Unfolded

	Volume	Open	Close
2019-01-01	1,000,000	\$150.00	\$155.00
2019-01-02	2,000,000	\$160.00	\$165.00
2019-01-03	3,000,000	\$155.00	\$160.00





A problem with RNNs: Vanishing Gradients



A problem with ~~RNNs~~-NNs: Vanishing Gradients



A problem with ~~RNNs~~-NNs: Vanishing Gradients

- When the gradients are too small - when you multiply a bunch of small gradients together you just end up with smaller and smaller gradients as you backpropagate
- The weights and biases of the earlier layers will not be effectively updated and can lead to an inaccurate model



A solution for NNs: ReLU

- One of the reasons ReLU is so effective is that it doesn't cause a small derivative
 - Convince yourself of this: what can the derivative of the ReLU function be?

A solution for RNNs: LSTMs and GRUs

- Long Short-Term Memory (LSTM)
- Gated Recurrent Unit (GRU)

Let's try it out!



